

# **Comp Photography (Sp 2015)**

## **HW 5**

Jonathan Hudgins

# Gradient Functions

## 1. imageGradientX

- a. The simplest and fastest way of accomplishing this is to use the numpy “roll” function to shift all pixels to the right (with wrapping).
- b. But because we are supposed to iterate over the pixels, I converted image to int, looped through y, then x, setting the result value to:  
`abs(image[y, x]-image[y,(x+1)%image.shape[1]])`
- c. Convert back to uint8 (no need to worry about overflow because the biggest absolute delta will be 255)

## 2. imageGradientY


- a. Same as “imageGradientX” only use:  
`abs(image[y, x]-image[(y+1)%image.shape[0],x])`

## 3. computeGradient


- a. Convert to float
- b. The simplest and fastest way is to use “scipy.signal.convolve2d”
- c. But because other libraries are disallowed I loop through y, then x (starting at 1 ending 1 less than image.shape)
- d. set result equal to sum of numpy.multiply of slice around y,x and the kernel

# Overview

Kernel 1  
7x7 Gaussian  
Threshold 0.05



Smoothed sigma 1  
Canny 15, 45





## Canny

First I smoothed the image using OpenCV “GaussianBlur” (7x7 with sigma of 1)

Then I used the OpenCV function “Canny” with threshold1 = 15 and threshold2 = 45.

Notice the single pixel width.



## Single Kernel

To create my own edges I tried a variety of kernels convolving a differential and gaussian kernel and then applying.

The differentials I tried were:

0 -1	-1 -1	-1 -2
-1 2	-1 3	-2 5

(with the other rows and columns 0)

The gaussians ranged from 3x3, 5x5, 7x7 with sigma 0.5, 1, 2.

I then convolved the kernel with the image. Then use absolute value of difference between pixel and median (distance from median). Normalize these values based on the maximum. Then threshold to 0 or 255. (Some nifty numpy matrix functions make all of this expedient!)

This image is pretty close (biggest difference is the edge width)